

# PROJECT : Building user-based recommendation model for Amazon.

## Description :

The dataset provided contains movie reviews given by Amazon customers. Reviews were given between May 1996 and July 2014.

## Data Dictionary :

UserID – 4848 customers who provided a rating for each movie  
Movie 1 to Movie 206 – 206 movies for which ratings are provided by 4848 distinct users

## Data Considerations

- All the users have not watched all the movies and therefore, all movies are not rated. These missing values are represented by NA.
- Ratings are on a scale of -1 to 10 where -1 is the least rating and 10 is the best.

## Analysis Task

- Exploratory Data Analysis:
  1. Which movies have maximum views/ratings?
  2. What is the average rating for each movie? Define the top 5 movies with the maximum ratings.
  3. Define the top 5 movies with the least audience.
    - Recommendation Model: Some of the movies hadn't been watched and therefore, are not rated by the users.
    - Netflix would like to take this as an opportunity and build a machine learning recommendation algorithm which provides the ratings for each of the users.
  4. Divide the data into training and test data
  5. Build a recommendation model on training data
  6. Make predictions on the test data

# 1.Import Necessary Libraries

```
In [11]: import numpy as np
import pandas as pd
import re
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
#import surprise
```

## 2.Load Dataset

```
In [12]: df=pd.read_csv("Amazon - Movies and TV Ratings.csv")
```

## 3.Explore Dataset

```
In [13]: df.head() # Check top 5 record
```

Out[13]:

	user_id	Movie1	Movie2	Movie3	Movie4	Movie5	Movie6	Movie7	Movie8	Movie9	...	Movie197	Movie198	Movie199	!
0	A3R5OBKS7OM2IR	5.0	5.0	NaN	NaN	NaN	NaN	NaN	NaN	NaN	...	NaN	NaN	NaN	
1	AH3QC2PC1VTGP	NaN	NaN	2.0	NaN	NaN	NaN	NaN	NaN	NaN	...	NaN	NaN	NaN	
2	A3LKP6WPMP9UKX	NaN	NaN	NaN	5.0	NaN	NaN	NaN	NaN	NaN	...	NaN	NaN	NaN	
3	AVIY68KEPQ5ZD	NaN	NaN	NaN	5.0	NaN	NaN	NaN	NaN	NaN	...	NaN	NaN	NaN	
4	A1CV1WROP5KTTW	NaN	NaN	NaN	NaN	5.0	NaN	NaN	NaN	NaN	...	NaN	NaN	NaN	

5 rows × 207 columns



```
In [14]: df.shape # printing No of rows and columns
```

Out[14]: (4848, 207)

```
In [15]: df_original=df.copy() # Make a copy of original dataset
```

```
In [16]: df.describe().T # Calculating some statistical data
```

Out[16]:

	count	mean	std	min	25%	50%	75%	max
<b>Movie1</b>	1.0	5.000000	NaN	5.0	5.00	5.0	5.0	5.0
<b>Movie2</b>	1.0	5.000000	NaN	5.0	5.00	5.0	5.0	5.0
<b>Movie3</b>	1.0	2.000000	NaN	2.0	2.00	2.0	2.0	2.0
<b>Movie4</b>	2.0	5.000000	0.000000	5.0	5.00	5.0	5.0	5.0
<b>Movie5</b>	29.0	4.103448	1.496301	1.0	4.00	5.0	5.0	5.0
...	...	...	...	...	...	...	...	...
<b>Movie202</b>	6.0	4.333333	1.632993	1.0	5.00	5.0	5.0	5.0
<b>Movie203</b>	1.0	3.000000	NaN	3.0	3.00	3.0	3.0	3.0
<b>Movie204</b>	8.0	4.375000	1.407886	1.0	4.75	5.0	5.0	5.0
<b>Movie205</b>	35.0	4.628571	0.910259	1.0	5.00	5.0	5.0	5.0
<b>Movie206</b>	13.0	4.923077	0.277350	4.0	5.00	5.0	5.0	5.0

206 rows × 8 columns

## Task 1: Which movies have maximum views/ratings?

```
In [17]: # Movie with highest views
```

```
df.describe().T['count'].sort_values(ascending=False)[:1].to_frame()
```

Out[17]:

	count
<b>Movie127</b>	2313.0

```
In [18]: # Movie with highest rating
df.drop('user_id',axis=1).sum().sort_values(ascending=False)[:1].to_frame()
```

Out[18]:

	0
Movie127	9511.0

**Task 2: What is the average rating for each movie? Define the top 5 movies with the maximum ratings.**

```
In [19]: df.drop('user_id',axis=1).mean().sort_values(ascending=False)[:5].to_frame()
```

Out[19]:

	0
Movie1	5.0
Movie66	5.0
Movie76	5.0
Movie75	5.0
Movie74	5.0

**Task 3: Define the top 5 movies with the least audience**

```
In [20]: df.drop('user_id',axis=1).mean().sort_values(ascending=True)[:5].to_frame()
```

Out[20]:

	0
Movie144	1.0
Movie67	1.0
Movie45	1.0
Movie58	1.0
Movie60	1.0

## Task 4 : Recommendation Model

```
In [21]: from surprise import Reader
from surprise import accuracy
from surprise import Dataset
from surprise.model_selection import train_test_split
from surprise import SVD
from surprise.model_selection import cross_validate
```

```
In [23]: df_melt=df.melt(id_vars=df.columns[0],value_vars=df.columns[1:],var_name="Movies",value_name="Rating")
```

```
In [24]: df_melt
```

```
Out[24]:
```

	user_id	Movies	Rating
0	A3R5OBKS7OM2IR	Movie1	5.0
1	AH3QC2PC1VTGP	Movie1	NaN
2	A3LKP6WPMP9UKX	Movie1	NaN
3	AVIY68KEPQ5ZD	Movie1	NaN
4	A1CV1WROP5KTTW	Movie1	NaN
...	...	...	...
998683	A1IMQ9WMFYKWH5	Movie206	5.0
998684	A1KLIKPUF5E88I	Movie206	5.0
998685	A5HG6WFZLO10D	Movie206	5.0
998686	A3UU690TWXCG1X	Movie206	5.0
998687	AI4J762YI6S06	Movie206	5.0

998688 rows × 3 columns

```
In [25]: rd=Reader()  
data=Dataset.load_from_df(df_melt.fillna(0),reader=rd)  
data
```

```
Out[25]: <surprise.dataset.DatasetAutoFolds at 0x2347b193220>
```

```
In [26]: trainset,testset=train_test_split(data,test_size=0.25)
```

```
In [27]: #Using SVD (Singular Value Decomposition)  
svd=SVD()  
svd.fit(trainset)
```

```
Out[27]: <surprise.prediction_algorithms.matrix_factorization.SVD at 0x23400fdffa0>
```

```
In [29]: pred=svd.test(testset)
```

```
In [30]: accuracy.rmse(pred)
```

RMSE: 1.0257

```
Out[30]: 1.025728887899078
```

```
In [31]: accuracy.mae(pred)
```

MAE: 1.0119

```
Out[31]: 1.0118749492558399
```

```
In [33]: cross_validate(svd,data,measures=['RMSE','MAE'],cv=3, verbose= True)
```

Evaluating RMSE, MAE of algorithm SVD on 3 split(s).

	Fold 1	Fold 2	Fold 3	Mean	Std
RMSE (testset)	1.0250	1.0268	1.0266	1.0261	0.0008
MAE (testset)	1.0116	1.0124	1.0122	1.0120	0.0003
Fit time	44.05	50.15	48.29	47.49	2.55
Test time	4.22	4.63	3.50	4.12	0.47

```
Out[33]: {'test_rmse': array([1.02501788, 1.02680059, 1.02659326]),
          'test_mae': array([1.01156378, 1.01239774, 1.01215008]),
          'fit_time': (44.04900121688843, 50.14578437805176, 48.28626322746277),
          'test_time': (4.220624208450317, 4.626606225967407, 3.501664876937866)}
```

```
In [34]: def repeat(ml_type,dframe):
          rd=Reader()
          data=Dataset.load_from_df(dframe,reader=rd)
          print(cross_validate(ml_type,data,measures=['RMSE','MAE'],cv=3,verbose=True))
          print("--"*15)
          usr_id = 'A3R50BKS70M2IR'
          mv = 'Movie1'
          r_u = 5.0
          print(ml_type.predict(usr_id,mv,r_ui = r_u,verbose=True))
          print("--"*15)
```

```
In [35]: repeat(SVD(),df_melt.fillna(df_melt['Rating'].mean()))
#repeat(SVD(),df_melt.fillna(df_melt['Rating'].median()))
```

Evaluating RMSE, MAE of algorithm SVD on 3 split(s).

	Fold 1	Fold 2	Fold 3	Mean	Std
RMSE (testset)	0.0856	0.0845	0.0881	0.0860	0.0015
MAE (testset)	0.0098	0.0097	0.0098	0.0098	0.0000
Fit time	47.95	47.77	49.64	48.46	0.84
Test time	3.75	5.13	4.14	4.34	0.58

```
{'test_rmse': array([0.08556113, 0.08450821, 0.08807983]), 'test_mae': array([0.00982059, 0.00974776, 0.00976051]), 'fit_time': (47.950905323028564, 47.77400255203247, 49.64251399040222), 'test_time': (3.752955198287964, 5.131466865539551, 4.141684293746948)}
```

```
-----
user: A3R5OBKS7OM2IR item: Movie1      r_ui = 5.00    est = 4.40    {'was_impossible': False}
user: A3R5OBKS7OM2IR item: Movie1      r_ui = 5.00    est = 4.40    {'was_impossible': False}
-----
```

```
In [36]: #trying grid search and find optimum hyperparameter value for n_factors
from surprise.model_selection import GridSearchCV
```

```
In [37]: param_grid = {'n_epochs':[20,30],
                        'lr_all':[0.005,0.001],
                        'n_factors':[50,100]}
```

```
In [38]: gs = GridSearchCV(SVD,param_grid,measures=['rmse','mae'],cv=3)
data1 = Dataset.load_from_df(df_melt.fillna(df_melt['Rating'].mean()),reader=rd)
gs.fit(data1)
```

```
In [39]: gs.best_score
```

```
Out[39]: {'rmse': 0.08478880296660705, 'mae': 0.009003663059166988}
```



```
In [40]: print(gs.best_score["rmse"])  
         print(gs.best_params["rmse"])
```

```
0.08478880296660705  
{'n_epochs': 30, 'lr_all': 0.001, 'n_factors': 50}
```

```
In [ ]:
```